

SPECIALIZED PROJECT REPORT

STUDYING AND DEVELOPING: DISTRIBUTED BARRIER ALGORITHMS USING THE HYBRID PROGRAMMING MODEL COMBINING MPI-3 AND C++11

Phạm Võ Quang Minh - 2111762

Vietnam National University - Ho Chi Minh City University of Technology

2025-01-03

Outline

1. Introduction

1.1 Motivation

1.2 Objectives

1.3 Scope

2. What's New?

2.1 MPI-3

2.2 C++11

3. Related Works

3.1 The Paper

3.2 Barrier Algorithm Selection

4. Algorithm & Simple Implementaion

4.1 Brook 2 process algorithm

4.2 Implementation using RMA

Operation

4.3 Preliminary Result

5. Conclusions

5.1 Accomplishments

5.2 Challenges and Learnings

5.3 Future Works

Outline

1. Introduction

1.1 Motivation

1.2 Objectives

1.3 Scope

2. What's New?

2.1 MPI-3

2.2 C++11

3. Related Works

3.1 The Paper

3.2 Barrier Algorithm Selection

4. Algorithm & Simple Implementaion

4.1 Brook 2 process algorithm

4.2 Implementation using RMA

Operation

4.3 Preliminary Result

5. Conclusions

5.1 Accomplishments

5.2 Challenges and Learnings

5.3 Future Works

1.1 Motivation

1.1.1 HPC and its Applications




Figure 1: Weather Simulation ^[1]

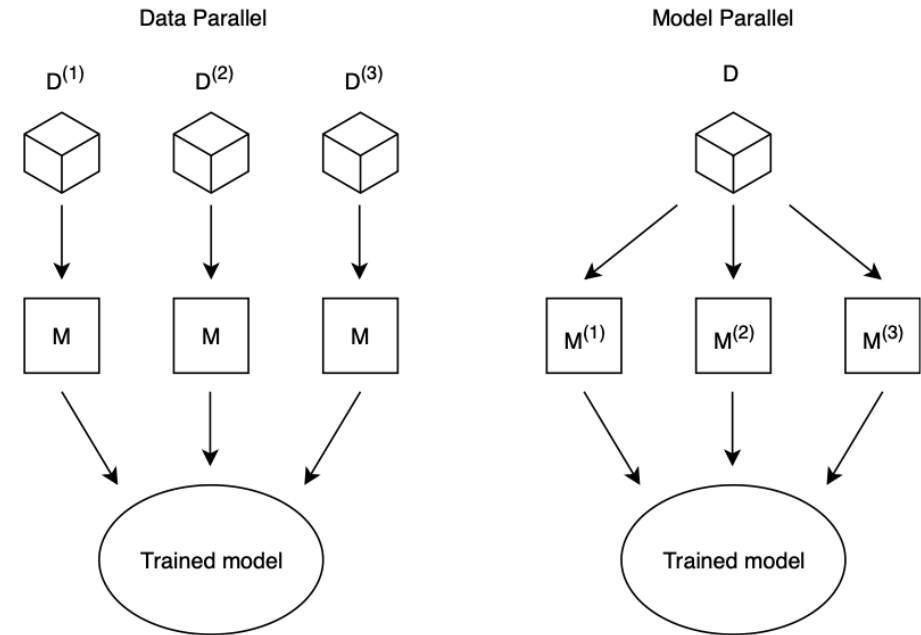


Figure 2: Distributed Machine Learning ^[2]

^[1]“NASA Global Weather Forecasting.” Accessed: Jan. 01, 2025. [Online]. Available: <https://www.nccs.nasa.gov/sci-tech/case-studies/nasa-global-weather-forecasting>


^[2]“A Survey on Distributed Machine Learning,” *ACM Computing Surveys*, vol. 53, doi: [10.1145/3377454](https://doi.org/10.1145/3377454).

1.1 Motivation

Moore's Law – The number of transistors on integrated circuit chips (1971-2018)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are linked to Moore's law.

Our World
in Data



Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)
The data visualization is available at [OurWorldInData.org](https://www.ourworldindata.org). There you find more visualizations and research on this topic.

Licensed under CC-BY-SA by the author Max Roser.

Figure 3: Moore's Law ^[1]

^[1]“Moore's Law Transistor Count 1971-2018.” Accessed: Jan. 01, 2025. [Online]. Available: https://upload.wikimedia.org/wikipedia/commons/8/8b/Moore%27s_Law_Transistor_Count_1971-2018.png

1.1 Motivation

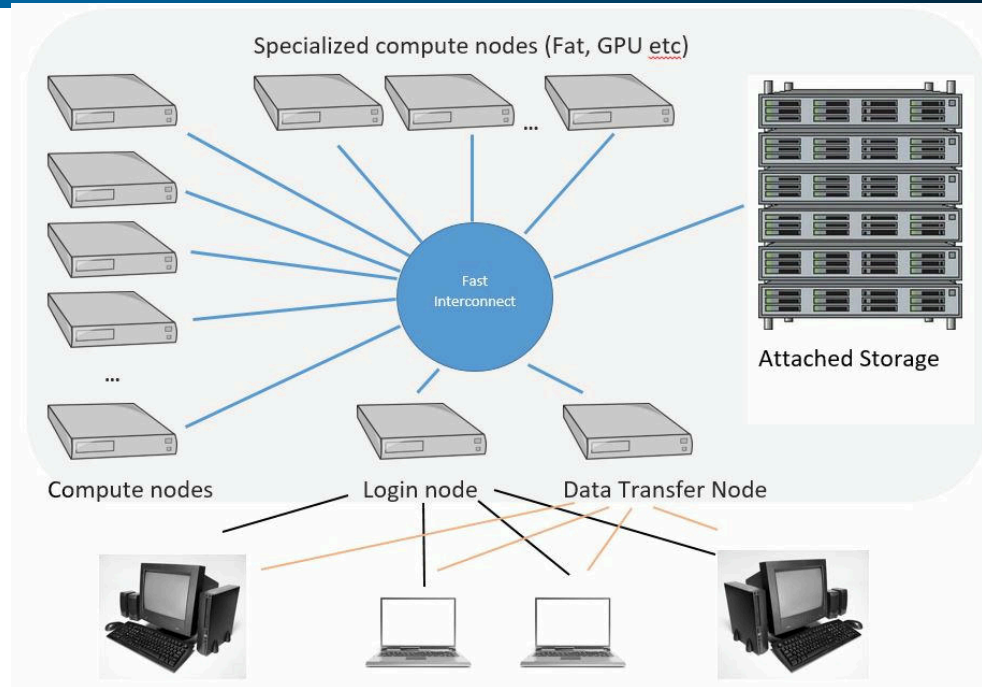


Figure 4: Multiple Computing Nodes connect to each other to form a HPC Cluster ^[1]

- Nodes are connected via a very fast network (Infiniband, Ethernet, etc.)

^[1]“What is an HPC Cluster?.” Accessed: Jan. 01, 2025. [Online]. Available: <https://www.hpc.iastate.edu/guides/introduction-to-hpc-clusters/what-is-an-hpc-cluster>

1.1 Motivation

What is a Barrier Algorithm?




Figure 5: Barrier Algorithm

1.1 Motivation

1.1.2 The Paper

- Quaranta et al ^[1] proposed to combine MPI-3 and C++11 (hybrid model)
 - Only implements a simple barrier algorithm using the hybrid model
- Implement and benchmark more complex barrier algorithms using the hybrid model

^[1]L. Quaranta and L. Maddegadara, “A Novel MPI+MPI Hybrid Approach Combining MPI-3 Shared Memory Windows and C11/C++11 Memory Model,” *Journal of Parallel and Distributed Computing*, vol. 157, pp. 125–144, Nov. 2021, doi: [10.1016/j.jpdc.2021.06.008](https://doi.org/10.1016/j.jpdc.2021.06.008).

1.2 Objectives

- Survey existing barrier synchronization algorithms
- Implements variants of barrier synchronization algorithms in the hybrid programming model.
- **Compare the performance of the implemented barrier algorithms with existing algorithms.**
- Propose methods to deploy selected barrier algorithms on the current HPC system.

1.3 Scope

- Explore key concepts in high-performance computing (HPC).
- Study MPI-3 and its programming model.
- Investigate C++11 threading and concurrency features.
- Research the hybrid programming model combining MPI-3 and C++11.
- Survey existing barrier synchronization algorithms.
- Implement a simple barrier synchronization algorithm with the hybrid model.

Outline

1. Introduction

1.1 Motivation

1.2 Objectives

1.3 Scope

2. What's New?

2.1 MPI-3

2.2 C++11

3. Related Works

3.1 The Paper

3.2 Barrier Algorithm Selection

4. Algorithm & Simple Implementaion

4.1 Brook 2 process algorithm

4.2 Implementation using RMA

Operation

4.3 Preliminary Result

5. Conclusions

5.1 Accomplishments

5.2 Challenges and Learnings

5.3 Future Works

2.1 MPI-3




Figure 6: Live program communicating with each other using MPI ^[1]

^[1]“Introduction to MPI - MPI Send and Receive.” Accessed: Jan. 01, 2025. [Online]. Available: https://pdc-support.github.io/introduction-to-mpi/03-mpi_send_recv/index.html

2.1 MPI-3

Traditional Message Passing

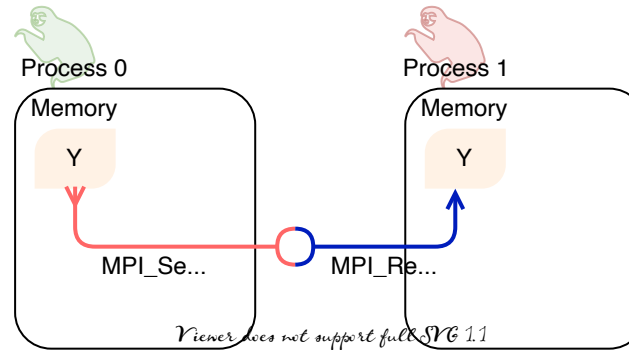


Figure 7: Point-to-point communication^[1]

- Point-to-point
- Explicit send and receive

^[1]“Introduction to MPI - MPI Send and Receive.” Accessed: Jan. 01, 2025. [Online]. Available: https://pdc-support.github.io/introduction-to-mpi/03-mpi_send_rcv/index.html

2.1 MPI-3

One-sided Communication




Figure 8: One-sided Communication^[2]

- Remote Memory Access
- Handshake is implicit

2.1.1 One-sided Communication

- Introduced in MPI-2
- Share mechanism:
 - Declare a window of memory to be shared
 - read/write without explicit send/receive

^[2]“Introduction to MPI - MPI Send and Receive.” Accessed: Jan. 01, 2025. [Online]. Available: https://pdc-support.github.io/introduction-to-mpi/03-mpi_send_rcv/index.html


2.1 MPI-3

- Simple operations:
 - MPI_Put
 - MPI_Get
 - MPI_Accumulate
- Atomic operations:
 - MPI_Get_accumulate
 - MPI_Fetch_and_op
 - MPI_Compare_and_swap

2.1 MPI-3

2.1.2 New Features in MPI-3

Separate Memory



Unified Memory




Figure 10: One-sided Communication^[1]

^[1]“Introduction to MPI - MPI Send and Receive.” Accessed: Jan. 01, 2025. [Online]. Available: https://pdc-support.github.io/introduction-to-mpi/03-mpi_send_rcv/index.html

2.2 C++11

Supports Multithreading

→ Allows Shared memory programming model instead of Distributed memory programming model

in the STL

→ No reliance on 3rd party libraries for multithreading programming

2.2 C++11

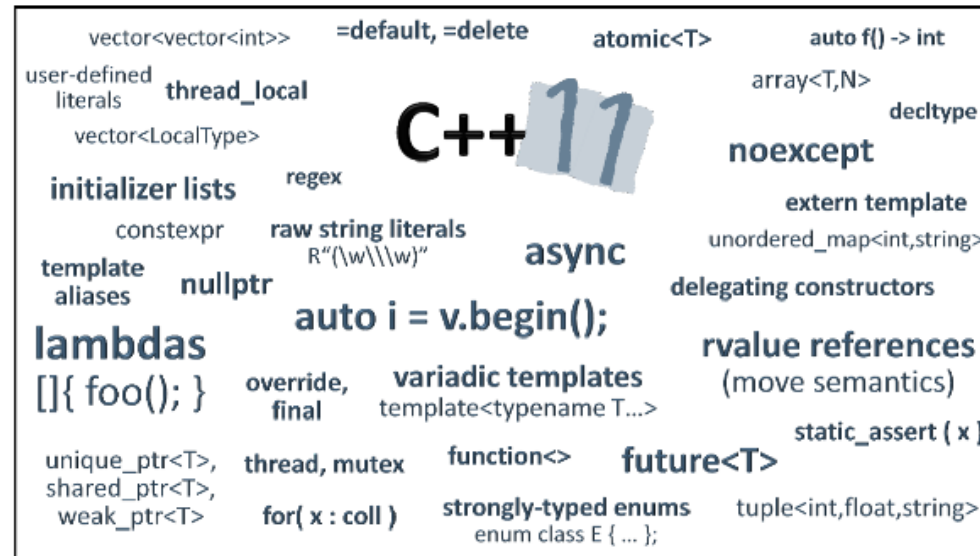


Figure 11: Features introduced in C++11 ^[1]

^[1]“Technical CPP Blog.” Accessed: Jan. 01, 2025. [Online]. Available: <https://www.cyberplusindia.com/blog/index.php/category/technical/cpp/>

Outline

1. Introduction

1.1 Motivation

1.2 Objectives

1.3 Scope

2. What's New?

2.1 MPI-3

2.2 C++11

3. Related Works

3.1 The Paper

3.2 Barrier Algorithm Selection

4. Algorithm & Simple Implementaion

4.1 Brook 2 process algorithm

4.2 Implementation using RMA

Operation

4.3 Preliminary Result

5. Conclusions

5.1 Accomplishments

5.2 Challenges and Learnings

5.3 Future Works

3.1 The Paper

- Quaranta et al^[1] proposed a hybrid model of MPI-3 and C++11
- Showed potential of combining MPI-3 and C++11 for a synchronization implementation
- Potential reduction in communication overhead

^[1]L. Quaranta and L. Maddegadara, “A Novel MPI+MPI Hybrid Approach Combining MPI-3 Shared Memory Windows and C11/C++11 Memory Model,” *Journal of Parallel and Distributed Computing*, vol. 157, pp. 125–144, Nov. 2021, doi: [10.1016/j.jpdc.2021.06.008](https://doi.org/10.1016/j.jpdc.2021.06.008).

3.2 Barrier Algorithm Selection

2 schools of thoughts:

- Linear
- Butterfly

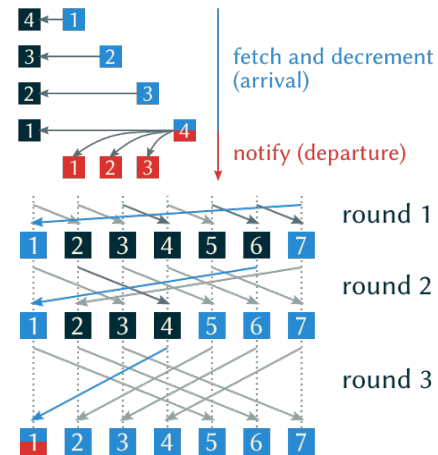


Figure 12: Barrier Algorithm Selection

Outline

1. Introduction

1.1 Motivation

1.2 Objectives

1.3 Scope

2. What's New?

2.1 MPI-3

2.2 C++11

3. Related Works

3.1 The Paper

3.2 Barrier Algorithm Selection

4. Algorithm & Simple Implementaion

4.1 Brook 2 process algorithm

4.2 Implementation using RMA

Operation

4.3 Preliminary Result

5. Conclusions

5.1 Accomplishments

5.2 Challenges and Learnings

5.3 Future Works

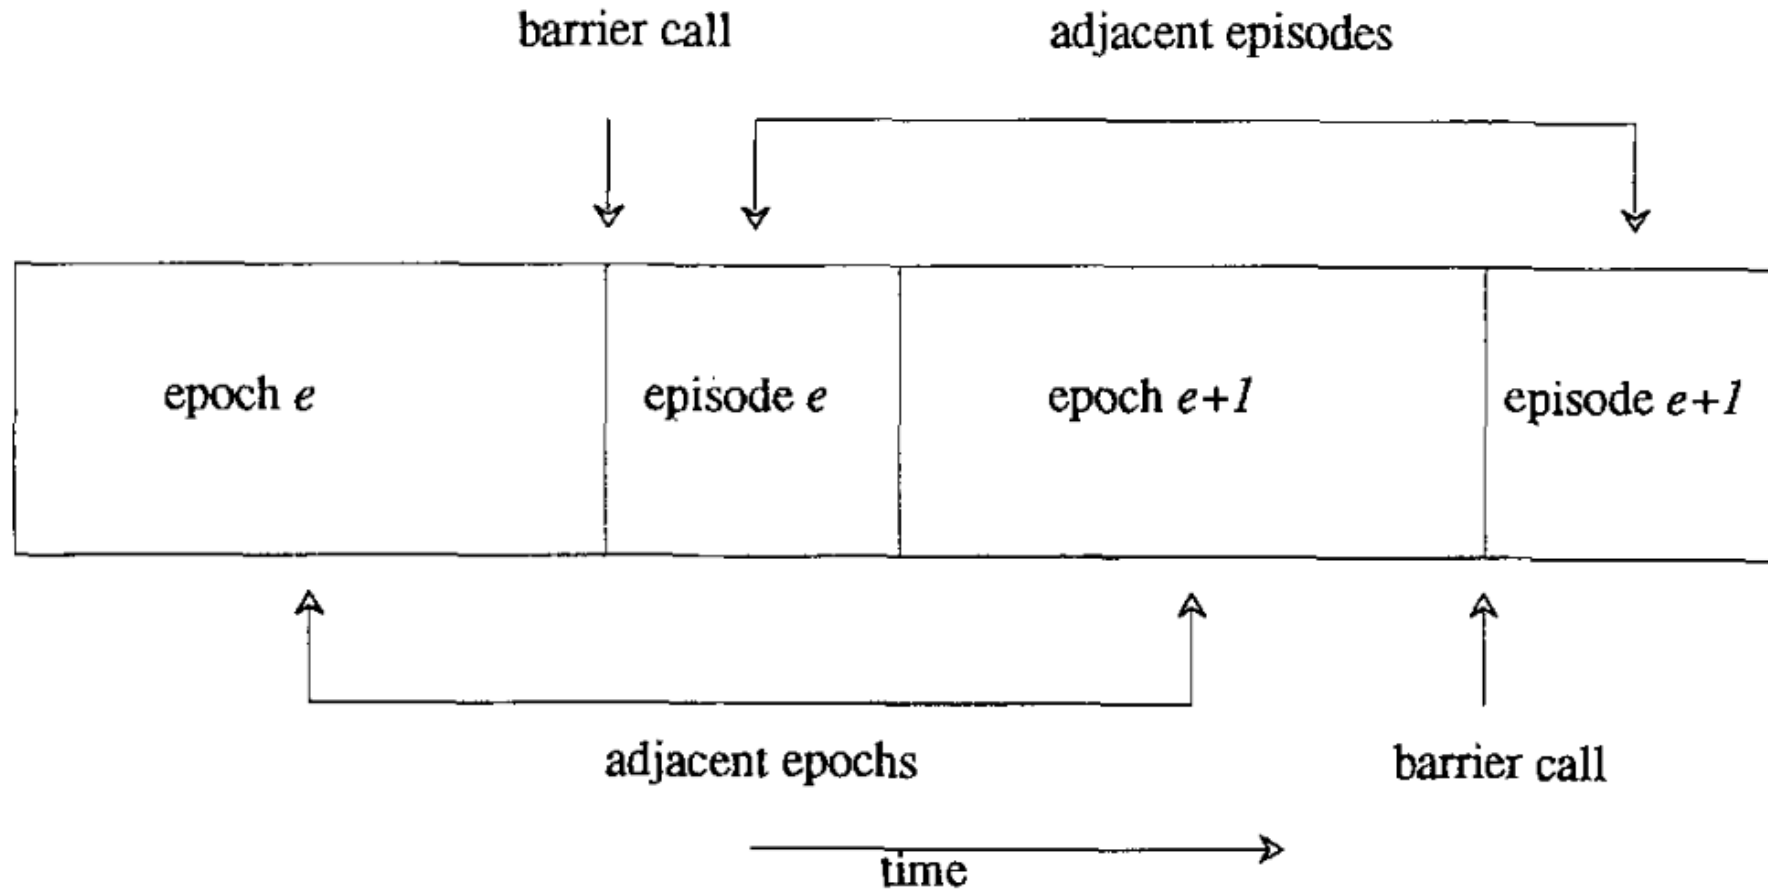
4.1 Brook 2 process algorithm

- The basis of many butterfly barrier algorithms
- Barrier for 2 threads of execution

BROOK-BARRIER-ALGORITHM

```
1 procedure Brook's-Barrier(SetByMyProcess, SetByTargetProcess)
2   while SetByMyProcess is true do wait
3   SetByMyProcess ← true
4   while SetByTargetProcess is false do wait
5   SetByTargetProcess ← false
6 end procedure
```

4.1 Brook 2 process algorithm



4.1 Brook 2 process algorithm

Step

2

3

Process 1

SetByProcess1 := true;

while not SetByProcess2 do wait;

Process 2

SetByProcess2 := true;

while not SetByProcess1 do wait;

Step

1

2

3

4

Process 1

while SetByProcess1 do wait;

SetByProcess1 := true;

while not SetByProcess2 do wait;

SetByProcess2 := false;

Process 2

while SetByProcess2 do wait;

SetByProcess2 := true;

while not SetByProcess1 do wait;

SetByProcess1 := false;

4.2 Implementation using RMA Operation

DISTRIBUTED-MEM-BROOK-BARRIER

```
1 procedure Brook's-Barrier(exposed_flag, win, target_rank)
2   while exposed_flag is true do wait //step 1
3   exposed_flag ← true //step 2
4   //step 3
5   while target_flag is false do wait
6     MPI_WIN_LOCK(win)
7     MPI_GET_ACCUMULATE(&target_flag, 0, BOOL, &target_flag, 1, BOOL,
      target_rank, 0, 1, BOOL, MPI_NO_OP, win);
8     MPI_WIN_FLUSH(win)
9     MPI_WIN_UNLOCK(win)
```

4.2 Implementation using RMA Operation

DISTRIBUTED-MEM-BROOK-BARRIER

```
10  end while
11  //step 4
12  false_value ← false
13  MPI_WIN_LOCK(win)
14  MPI_ACCUMULATE(&>false_value, 1, BOOL, target_rank, 0, 1, BOOL,
    MPI_REPLACE, win);
15  MPI_WIN_FLUSH(win)
16  MPI_WIN_UNLOCK(win)
17  end procedure
```

4.3 Preliminary Result

4.3.1 Test Environment

- Hardware: Apple M1 chip
- Compiler: MPICH's mpic++ compiler
- MPI Implementation: MPICH
- Compilation flags: -std=c++11
- Operating System: macOS

```
int main(int argc, char **argv) {  
    // init the mpi world  
    MPI_Init(&argc, &argv);  
    MPI_Comm comm = MPI_COMM_WORLD;  
  
    // get number of ranks  
    int n_ranks;  
    MPI_Comm_size(comm, &n_ranks);
```

4.3 Preliminary Result

```
// get current rank
int rank;
MPI_Comm_rank(comm, &rank);

if (n_ranks != 2) {
    if (rank == 0) {
        fprintf(stderr, "This program requires exactly 2 processes.\n");
    }
    MPI_Abort(comm, 1);
}

// set the target
int target_rank = 1 - rank;

// create a window
bool exposed_bool{false};
```

4.3 Preliminary Result

```
MPI_Win win_buffer_handler;  
MPI_Win_create(&exposed_bool, sizeof(bool), sizeof(bool), MPI_INFO_NULL,  
comm,  
                &win_buffer_handler);  
  
barrier_brook(exposed_bool, win_buffer_handler, target_rank);  
barrier_brook(exposed_bool, win_buffer_handler, target_rank);  
barrier_brook(exposed_bool, win_buffer_handler, target_rank);  
barrier_brook(exposed_bool, win_buffer_handler, target_rank);  
  
printf("Process %d: reached destination\n", rank);  
  
return MPI_Finalize();  
}
```

4.3 Preliminary Result

4.3.2 Compilation and Execution

The program was compiled using the following command:

```
mpic++ -std=c++11 src/brook.cpp -o brook
```

Execution was performed using MPICH's mpirun with two processes:

```
mpirun -np 2 ./brook
```


Outline

1. Introduction

1.1 Motivation

1.2 Objectives

1.3 Scope

2. What's New?

2.1 MPI-3

2.2 C++11

3. Related Works

3.1 The Paper

3.2 Barrier Algorithm Selection

4. Algorithm & Simple Implementaion

4.1 Brook 2 process algorithm

4.2 Implementation using RMA

Operation

4.3 Preliminary Result

5. Conclusions

5.1 Accomplishments

5.2 Challenges and Learnings

5.3 Future Works

5.1 Accomplishments

- Studied and familiarized myself with MPI and its One-Sided Communication techniques
- Explored C++11 threading and concurrency features
- Successfully implemented the two-process Brooks barrier algorithm utilizing MPI's Remote Memory Access (RMA) Operations

5.2 Challenges and Learnings

5.2.1 Initial Challenges:

- Faced a steep learning curve with parallel programming concepts.
- Extensive effort required to comprehend the MPI standard documentation and its historical context.

5.2.2 Key Learnings:

- Gained a robust understanding of memory models and synchronization techniques.
- Developed insight into one-sided communication mechanics in MPI.

5.2.3 Adaptation Effort:

- Adapting Brook's barrier algorithm to distributed memory was straightforward but required a comprehensive system understanding.

5.2 Challenges and Learnings

5.2.4 Outcomes:

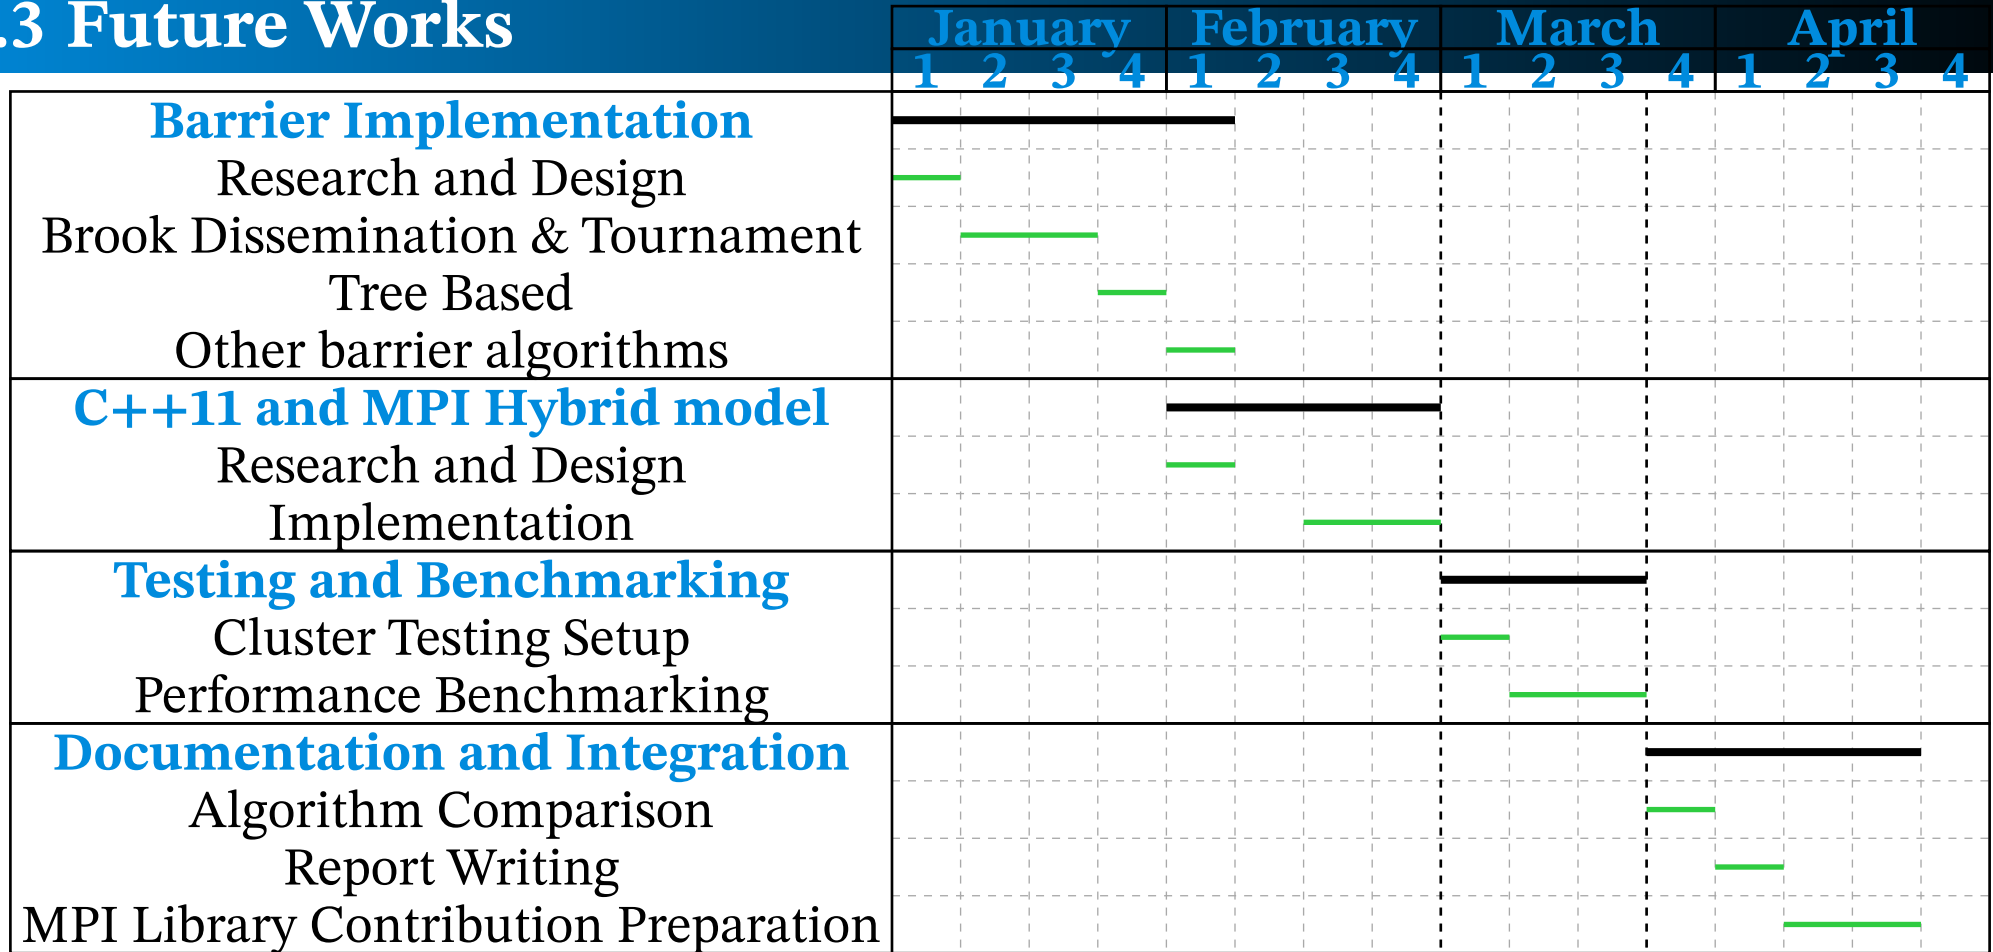
- Established a solid foundation in distributed memory systems and synchronization strategies.
- Successfully adapted the two-process barrier algorithm to distributed memory.
- Future Opportunities:
 - Yet to implement a fully functional barrier algorithm for hybrid memory models.
 - Room for further exploration and development in this domain.

5.3 Future Works

5.3.1 Plan

- Understanding of C++ and MPI Memory model to implement C++ and MPI Hybrid model
- Adapting and Implementing barrier algorithms using MPI's One-Sided Communication
- Testing of proposed algorithms on existing clusters
- Performance benchmarking

5.3 Future Works



Implementation

Testing